CMSC 180
Introduction to Parallel Computing
Second Semester AY 2024-2025

Laboratory Research Problem 05
Distributed Computation of the Mean Square Error of a
Column in a Matrix and a Vector

Introduction

The answer to question 5 of Laboratory Research Problem 04 should have been a one-to-many personalized broadcast (1MPB) of the matrix **X**. If you have implemented 1MPB in your `lab04`, it is expected that you will observe a significant improvement on average runtime as *n* and *t* increases. The 1MPB distributed parts of the matrix **X** to different slave processes which can be processed concurrently by different machines. But the respective Mean Square Errors (MSEs) of the matrix **X**'s columns with the vector **y** can not yet be computed by the slave processes without the full vector **y**. To fully compute in distributed manner the respective MSEs, the vector **y** needs to be replicated among the slave processes. The efficient way to do this is to distribute the vector **y** in a one-to-many broadcast (1MB) manner. The respective MSEs should be compiled in a vector **e** (recall Equation 1 in Laboratory Research Problem 01), the elements of which would have been computed independently by the slave processes and which should be sent back to the master in a many-to-one personalized reduction (M1PR) manner to be rebuilt into one big vector by the master process.

Research Activity 1: How to do it?

1. Extend your `lab04` into `lab05` which...
   (1) should include the distribution of the full vector **y** in a 1MB manner;
   (2) should compute the Mean Square Error into a vector **e** of the respective columns of the received parts of matrix **X** with the full vector **y**;
   (3) should communicate back to the master the respective parts of vector **e** in a M1PR manner.

2. As in `lab04`, the master process must take note of the `time_before` before distributing the matrix **X** and the vector **y** to the different slave processes and then take note of the `time_after` after the master has rebuilt the full vector **e**. The master process must report the `time_elapsed`.

3. Each of the slave processes must take note of the time_elapsed ONLY in computing the Mean Square Error. Strictly, this means that a slave process must take note of the `time_before` after receiving the parts of matrix **X** and vector **y** and before computing for parts of the vector **e** and then take note of the `time_after` after the parts of vector **e** has been fully computed and before the vector **e** is sent back to the master. The `time_elapsed` is the time spent by each slave process in actual computation. The slave must report the `time_elapsed`.

4. Conduct this activity using different machines, where the slave processes in each machine are running in a core-affine manner.

5. Fill in the **Table 1** below with your time readings from the master process only.

Table 1. Time elapsed as reported by the master process.					
n	*t*	**Time Elapsed (seconds)**			**Average Runtime (seconds)**
		Run 1	**Run 2**	**Run 3**	
20,000	2				
20,000	4				
20,000	8				
20,000	16				
25,000	2				
25,000	4				
25,000	8				
25,000	16				
30,000	2				
30,000	4				
30,000	8				
30,000	16				

6. Fill in **Table 2** below with your time readings from the slaves. Fill in only the maximum for each run.

Table 2. Time elapsed as reported by the slave processes.					
n	*t*	**Time Elapsed (seconds)**			**Average Runtime (seconds)**
		Max of Run 1	**Max of Run 2**	**Max of Run 3**	
20,000	2				
20,000	4				
20,000	8				
20,000	16				
25,000	2				
25,000	4				
25,000	8				
25,000	16				
30,000	2				
30,000	4				
30,000	8				
30,000	16				

Research Activity 2: Communication and Computation

The average runtime that you obtained in **Table 1** above includes the computational cost (in terms of time) of both communication and computation of the whole distributed computation. Using any software (a spreadsheet will do), create a 3D plot of n, t and average runtime, where the average runtime is in the vertical axis. Name the plot as **Figure 1** in your report.

Alternatively, if you can not create a 3D plot, create a 2D graph of t and average runtime, where t is plotted as the horizontal axis and the average runtime is the vertical axis. You will be plotting three plots on the same 2D graph where each plot corresponds to each of the n.

Discuss the pattern that you observed for increasing n and t.

Research Activity 3: Computation Only

The average runtime that you compiled in **Table 2** above presents only the mean of the maximum computation cost spent by the slave processes. Why do we need to report the maximum of the `elapsed_time` for each run? Why not the average of each run?

As in Research Activity 2, plot n, t and average runtime. Name the plot as **Figure 2** in your report.

Discuss the pattern that you observed for increasing n and t.

Research Activity 4: Performance Metrics

For each problem size n, Compare the plots in Figures 1 and 2. You must be able to observe that the plots in Figure 1 has a higher runtime than the plots in Figure 2. Obviously because the runtime in Figure 1 includes both the communication and computation time while the runtime in Figure 2 only shows the computation time. But, how much did your `lab05` spent in communication over computation? Aside from communication and computation, do you think there are overhead that was spent by your `lab05`? If so, what are those (i.e., if there are <u>idling</u>, where are those in your code, if there are <u>excess computation</u>, why was the excess included)?

Using your result in Laboratory Research Problem 01 as serial runtime T_S (should be constant for each n), the average runtime you obtained in Table 1 as the parallel runtime T_P, and the number of processors p as t, report as **Table 3** (see next page) the parallel overhead T_O, the parallel speedup S, the parallel efficiency E, and the parallel cost pT_P for increasing n and t.

Separately discuss the patterns of performance metrics that you obtained for each n and t. Did you observe superlinearity? Is your implementation cost-optimal?

Reaction to possible problem in the activity:

Problem 1: *I was not able to implement in* `lab04` *and also here in* `lab05` *the 1MPB, 1MB and M1PR.*
Reaction 1: Include in your report that you were not able to but articulate in the report what would have been different had these broadcast schemes been implemented correctly.

n	t	Serial T_{S}	Parallel			
			T_{O}	S	E	pT_{P}
20,000	2					
20,000	4					
20,000	8					
20,000	16					
25,000	2					
25,000	4					
25,000	8					
25,000	16					
30,000	2					
30,000	4					
30,000	8					
30,000	16					

Table 3. Performance metrics of the parallel, distributed Pearson Correlation Coefficient.